

# Lecture 11: Computational Learning Theory

## PAC Learning

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.
- [Today] In Computational Learning Theory, we will treat the learning algorithm as a computational process. This requires making a few things explicit (Discuss in pairs what they can be):

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.
- [Today] In Computational Learning Theory, we will treat the learning algorithm as a computational process. This requires making a few things explicit (Discuss in pairs what they can be):
  - How does the algorithm “interact” with the distribution ?

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.
- [Today] In Computational Learning Theory, we will treat the learning algorithm as a computational process. This requires making a few things explicit (Discuss in pairs what they can be):
  - How does the algorithm “interact” with the distribution ?
  - What information does the algorithm have about the “learning problem” ?

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.
- [Today] In Computational Learning Theory, we will treat the learning algorithm as a computational process. This requires making a few things explicit (Discuss in pairs what they can be):
  - How does the algorithm “interact” with the distribution ?
  - What information does the algorithm have about the “learning problem” ?
  - How “long” is the algorithm allowed to run etc ?

# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.
- [Today] In Computational Learning Theory, we will treat the learning algorithm as a computational process. This requires making a few things explicit (Discuss in pairs what they can be):
  - How does the algorithm “interact” with the distribution ?
  - What information does the algorithm have about the “learning problem” ?
  - How “long” is the algorithm allowed to run etc ?
  - How will the output estimator be “represented” ?



# Computational Learning Theory

- So far our questions were of the form — Given  $m$  samples iid from a distribution  $\mathbb{P}_x$ , can we upper bound the error as  $\epsilon = O(m^{-\alpha})$ ?
- In particular, we have discussed how to use properties of the distribution to give better statistical “rates”.
- [Today] In Computational Learning Theory, we will treat the learning algorithm as a computational process. This requires making a few things explicit (Discuss in pairs what they can be):
  - How does the algorithm “interact” with the distribution ?
  - What information does the algorithm have about the “learning problem” ?
  - How “long” is the algorithm allowed to run etc ?
  - How will the output estimator be “represented” ?

Finally, this proposes an answer to the question what can be “learned” under various restrictions.

# Binary classification Setting

## Some terminologies

- **Instance space:**  $\mathcal{X}$  e.g.  $\mathbb{R}^2, \{0, 1\}^d, \mathbb{R}^d$  etc.
- **Label space:**  $\mathcal{Y} = +1, -1$
- **Hypothesis/Concept classes** are represented by  $\mathcal{C}, \mathcal{H}, \mathcal{F}$ . They are sets of maps from  $\mathcal{X}$  to  $\mathcal{Y}$ . (In other words, classes of labelling functions) E.g.
  - CONJUNCTIONS e.g.  $x_1 \wedge x_3 \wedge x_5$
  - DISJUNCTIONS e.g.  $x_2 \vee x_3 \vee x_5$
  - Linear halfspaces e.g.  $\sum_{i=1}^d w_i x_i \geq b$
- **Data Distribution**  $\mathbb{P}_x$  over  $\mathcal{X}$
- **Example Oracle:** An oracle  $\text{Ex}(c; \mathbb{P}_x)$  that samples  $x \sim \mathbb{P}_x$  and returns  $(x, c(x))$ .
- **Target Concept** Refer to  $c$  as the “target concept” (ground truth).

# Learning Algorithm

- **Learning algorithm** An algorithm  $\mathcal{A}$ 
  - for learning concept class  $\mathcal{C}$
  - with hypothesis class  $\mathcal{H}$
  - can call the example oracle  $\text{Ex}(c; \mathbb{P}_x)$  many times
  - and must return some  $h \in \mathcal{H}$ .
- Two sources of randomisation :
  - **Randomness from data** Inherently, due to the randomisation of  $\text{Ex}(c; \mathbb{P}_x)$ ,  $\mathcal{A}$  is always randomised. This randomness is from  $\mathbb{P}_x$ .
  - **Randomness from algorithm** After receiving data from  $\text{Ex}(c; \mathbb{P}_x)$ ,  $\mathcal{A}$  can flip an unbiased coin and introduce further randomness into the algorithm. Let the joint distribution over  $\mathbb{P}_x$  and internal coin flips of  $\mathcal{A}$  be  $\mathbb{P}$ .

# Probably Approximately Correct Learnability: Attempt 1

## Definition (PAC learning)

A concept class  $\mathcal{C}$  is PAC learnable with hypothesis class  $\mathcal{H}$  if there exists a learning algorithm  $\mathcal{A}$  such that for all distributions  $\mathbb{P}_x$ , concept  $c \in \mathcal{C}$ , and  $\epsilon, \delta > 0$ , if  $\mathcal{A}$  is given access to  $\text{Ex}(c; \mathbb{P}_x)$  and knows  $\epsilon, \delta$ ,  $\mathcal{A}$  returns  $h \in \mathcal{H}$  such that with probability at least  $1 - \delta$ , over inner randomisation of  $\text{Ex}(c; \mathbb{P}_x)$  and  $\mathcal{A}$  we have that  $\mathbb{P}_x[h(x) \neq c(x)] \leq \epsilon$ .

- If  $\mathcal{A}$  runs in time  $\text{poly}\left(\frac{1}{\epsilon}, \frac{1}{\delta}\right)$  then  $\mathcal{C}$  is **efficiently PAC learnable**.
- If  $\mathcal{C}$  is learnable with  $\mathcal{H} = \mathcal{C}$ , then we say  $\mathcal{C}$  is **proper learnable**. Otherwise, it is referred to as **improper learnable**. We will focus on proper PAC learnability for now.
- Number of times  $\mathcal{A}$  calls  $\text{Ex}(c; \mathbb{P}_x)$  is equal to the sample size  $m$ .
- So far, we have written  $\epsilon$  as function of  $m$  i.e.  $\epsilon(m, \delta)$  is the error rate.  $m$  is the sample size or the smallest possible  $m$  is the sample complexity.

# Understanding the definition

What are some things or questions that stand out to you about learnability in this definition ?

- **Efficiency**

- What is one unit of time ?
- What are possible reasons of inefficiency ?
- What kinds of computational constraints are required on  $h$  ?

- **Available information to  $\mathcal{A}$**

- What does  $\mathcal{A}$  know and what does  $\mathcal{A}$  not know ?
- What are some possible changes to  $\text{Ex}(c; \mathbb{P}_x)$  that can simulate real environments ? How can they change a class' learnability ?

**Discuss in pairs**

# Understanding the definition

- Efficiency.
  - What is one unit of time ?  
Call to  $\text{Ex}(c; \mathbb{P}_x)$  takes unit time. The algorithm is run on a turing machine.
  - What are possible reasons of inefficiency ?  
Exponential sample complexity or exponential running time.
  - What kinds of computational constraints are required on  $h$  ?  
 $h$  needs to be poly evaluable, otherwise trivial
- Available information
  - What does  $\mathcal{A}$  know and what does  $\mathcal{A}$  not know ?  
Knows  $\mathcal{C}$  but not  $c$ . Does not know  $\mathbb{P}_x$ .
  - What are some changes to  $\text{Ex}(c; \mathbb{P}_x)$  that can simulate real environments ?  
Noisy Oracle (RCN, Massart, Tsybakov), Positive/Negative only, Membership Query, Statistical Query
- It attempts to separate the two things
  - Having sufficient data
  - Being able to compute the estimator/hypothesis from the data

## Learning Axis-Aligned Rectangles

- Let  $\mathcal{X} = \mathbb{R}^2$ ,  $\mathcal{Y} = \{+1, 0\}$
- $\mathcal{C}$  is the class of Axis-Aligned Rectangle Classifiers. A concept  $c \in \mathcal{C}$  labels  $x \in \mathcal{X}$  positive (+1) if  $x$  lies inside the rectangle and 0 o.w.

### Theorem

*The concept class of axis aligned rectangles is efficiently proper PAC learnable.*

Proof:

- Algorithm  $\mathcal{A}$  chooses  $m = \frac{4}{\epsilon} \log \left( \frac{4}{\delta} \right)$ , queries  $\text{Ex}(c; \mathbb{P}_x)$   $m$  times and outputs the smallest axis-aligned rectangle  $R'$  that contains all +ve points.
- Let  $R$  be the target rectangle. Choose 4 regions  $T_1, T_2, T_3, T_4$  along the inner sides of  $R$  such that each region has mass  $\frac{\epsilon}{4}$  under  $\mathbb{P}_x$ . Note that if  $\text{Ex}(c; \mathbb{P}_x)$  returns at least one point in all of these regions with probability greater than  $1 - \delta$ , it suffices for us.
- Let  $A_i$  be the event that  $\text{Ex}(c; \mathbb{P}_x)$  upon  $m$  calls does not return any point in  $T_i$ . Show  $\mathbb{P}[\bigcup_i A_i] \leq 4 \exp \left( -\frac{m\epsilon}{4} \right)$
- Setting  $m = \frac{4}{\epsilon} \log \left( \frac{4}{\delta} \right)$  completes the proof.

## Probably Approximately Correct Learnability: Attempt 2

**Issue:** Previous definition does not account for the size of the concept class or the instance space.

- **Representation scheme for concept class:**  $\rho : (\Sigma \cup \mathbb{R})^* \rightarrow \mathcal{C}$  is a representation scheme for  $\mathcal{C}$ . e.g.  $\rho((x_1, y_1), (x_2, y_2)) =$  axis-aligned rectangle with bottom left corner at  $(x_1, y_1)$  and top right corner in  $(x_2, y_2)$ . (Unit cost to represent alphabets in  $\Sigma$  and numbers in  $\mathbb{R}$ )
- **Size of representations** The function  $\text{size} : (\Sigma \cup \mathbb{R})^* \rightarrow \mathbb{N}$  measures the size of a representation in  $(\Sigma \cup \mathbb{R})^*$ .
- **Size of concept:** A size of a concept is the minimum size over all representations in that representation scheme
$$\text{size}(c) = \min_{\sigma: \rho(\sigma)=c} .\text{size}(\sigma)$$

What are some examples where the choice of  $\rho$  affects the size of a concept?

- **Instance size:** Instances  $x \in \mathcal{X}$  also has an associated size e.g. memory to store. We denote  $\mathcal{X}_d$  as an instance space where all  $x \in \mathcal{X}_d$  has size  $d$ .

What are some examples of sizes of instance spaces?

Often these are clear from context but sometimes need further thought.



## Probably Approximately Correct Learnability: Attempt II

For  $d \geq 1$ , let  $\mathcal{C}_d$  be a concept class over  $\mathcal{X}_d$ . Consider instance space  $\mathcal{X} = \bigcup_{d=1}^{\infty} \mathcal{X}_d$  and the corresponding concept class  $\mathcal{C} = \bigcup_{d=1}^{\infty} \mathcal{C}_d$ .

### Definition (PAC learning)

A concept class  $\mathcal{C}$  is PAC learnable with hypothesis class  $\mathcal{H}$  if there exists a learning algorithm  $\mathcal{A}$  such that for all  $d > 0$ , all distributions  $\mathbb{P}_x$  over  $\mathcal{X}_d$ , concept  $c \in \mathcal{C}_d$ , and  $\epsilon, \delta > 0$ , if  $\mathcal{A}$  is given access to  $\text{Ex}(c; \mathbb{P}_x)$  and knows  $\epsilon, \delta, \text{size}(c)$ , and  $d$ ,  $\mathcal{A}$  returns  $h \in \mathcal{H}$  such that with probability at least  $1 - \delta$ , over inner randomisation of  $\text{Ex}(c; \mathbb{P}_x)$  and  $\mathcal{A}$  we have that  $\mathbb{P}_x[h(x) \neq c(x)] \leq \epsilon$ .

**Efficient PAC learnability:**  $\mathcal{A}$  should run in time polynomial in  $\frac{1}{\epsilon}, \frac{1}{\delta}, \text{size}(c)$ , and  $d$ . Usually  $\text{size}(c)$  is bounded by some polynomial in  $d$  and hence can be ignored.

# Learning CONJUNCTIONS

Now we will see an example of PAC Learning Attempt II

- Let  $\mathcal{X}_d = \{0, 1\}^d$ ,  $\mathcal{Y} = \{0, 1\}$
- $\text{CONJUNCTIONS}_d$  over  $d$  boolean variables  $z_1, \dots, z_d$ 
  - **literal** is a variable or its negation
  - **conjunction** is an AND of literals.
- A conjunction can be represented with two sets  $P, N \subseteq [d]$

$$c_{P,N} = \bigwedge_{i \in P} z_i \wedge \bigwedge_{j \in N} \bar{z}_j$$

- The class of  $\text{CONJUNCTIONS}_d$  is the set of all conjunctions.

$$\text{CONJUNCTIONS}_d = \{c_{P,N} \mid P, N \subseteq [d]\}$$

- Note an efficient representation scheme: size  $c_{P,N} \leq d$

## Theorem (learning conjunctions)

*The concept class  $\mathcal{C} = \bigcup_{d \geq 1} \text{CONJUNCTIONS}_d$  is efficiently PAC learnable.*

## Proof of learning conjunctions

**Proof** Let  $c^*$  be the target concept. The proof follows these three steps

- **Algorithm** Fix  $m \geq \frac{2d}{\epsilon} \log \left( \frac{2d}{\delta} \right)$  and run the following algorithm.
  - Start with  $c_{[d],[d]}$ ; call  $\text{Ex}(c^*; \mathcal{D})$   $m$  times.
  - For every +ve example, eliminate all literals from  $c_{[d],[d]}$  that cause it to be zero. Return the resultant conjunction  $h$ .
- **“Correct”** Call a literal  $\ell$  “bad” if  $\mathbb{P}_x[c^*(x) = 1 \wedge \ell(x) = 0] \geq \frac{\epsilon}{2d}$ .
  - Note by construction,  $\mathbb{P}_x[h(x) \neq c^*(x)] = \mathbb{P}_x[h(x) = 0 \wedge c^*(x) = 1]$ .
  - Let  $B$  be the set of bad literals and  $h$  contain no literals in  $B$ . Then,  
$$\mathbb{P}_x[h(x) = 0 \wedge c^*(x) = 1] \leq \sum_{\ell \in B} \mathbb{P}_x[h(x) = 0 \wedge \ell(x) = 1] \leq \epsilon$$
- **“Approximately”** Now, we need to prove that  $h$  contains no bad literals. Let  $A_\ell$  be the event that  $\ell$  is not eliminated by the algorithm after  $m$  calls.
  - Bound  $\mathbb{P}[A_\ell] \leq (1 - \frac{\epsilon}{2d})^m \leq \exp(-\frac{\epsilon m}{2d})$ .
  - $\mathbb{P}[\text{at least 1 “bad” literal remain}] \leq \mathbb{P}\left[\bigcup_{\ell \in B} A_\ell\right] \leq \sum_{\ell=0}^{2d} \exp(-\frac{\epsilon m}{2d})$
  - Use  $m \geq \frac{2d}{\epsilon} \log \left( \frac{2d}{\delta} \right)$  to show that all bad literals are eliminated with probability  $1 - \delta$ .

## Exercise: Learning k-CNF

A different approach to proving learnability — **By Reduction**

- Let  $\mathcal{X}_d = \{0, 1\}^d$ ,  $\mathcal{Y} = \{0, 1\}$
- $k\text{-CNF}_d$  over  $d$  boolean variables.
  - Set of  $k$ -tuples  $S = \{S^1, \dots, S^p\}$ .  $\forall i: S^i \subset [d]$  and  $|S^i| = k$ .
  - $c_S(x) = \bigwedge_{i=1}^p (\bigvee_{j=1}^k x_{S_j^i})$

### Theorem (learning k-CNF)

*The concept class  $\mathcal{C} = \bigcup_{d \geq 1} k\text{-CNF}_d$  is efficiently PAC learnable.*

**Proof by Reduction** We need the concept of *monotone conjunctions*.  
*Monotone conjunctions* are conjunctions without negated literals.

- **Step 1: (Reduction)** Show these exist: instance space  $Z_k$ , a map  $\phi: \{0, 1\}^d \rightarrow Z_k$ , and a bijection  $\varphi$  between  $k\text{-CNF}_d$  and monotone conjunctions on  $Z_k$ .
- **Step 2: (Algorithm)** Show that the algorithm for learning conjunctions also learns monotone conjunctions. Hence,  $\mathcal{C}$  is PAC learnable.
- **Step 3: (Reduction is polynomial)** Show that  $\phi, \phi^{-1}$ , and  $\varphi$  are polynomially evaluable and thus  $\mathcal{C}$  is efficiently PAC learnable.